

KUBERNETES

AN INTRODUCTION

MUST USE CONTAINERS



KUBERNETES:

Kubernetes is a portable, extensible, open source platform for managing containerized workloads and services, that facilitates both declarative configuration and automation. It has a large, rapidly growing ecosystem. Kubernetes services, support, and tools are widely available.

Ex. KUBERNETES DOES NOT HAVE THE CONCEPT OF NETWORKING. FOR IT IS A PLUG IN TO MOUNT ON TOP. AS STORAGE AND OTHER THINGS

WHAT DOES "KUBERNETES" MEAN?

Greek for “pilot” or
“Helmsman of a ship”

THE IDEA IS THAT IT DRIVES YOUR
INFRASTRUCTURE



EX: IT'S AN ORCHESTRA WHERE EVERY MUSICIAN IS A COMPONENT OR A RESOURCE, AND KUBERNETS IS THE DIRECTOR

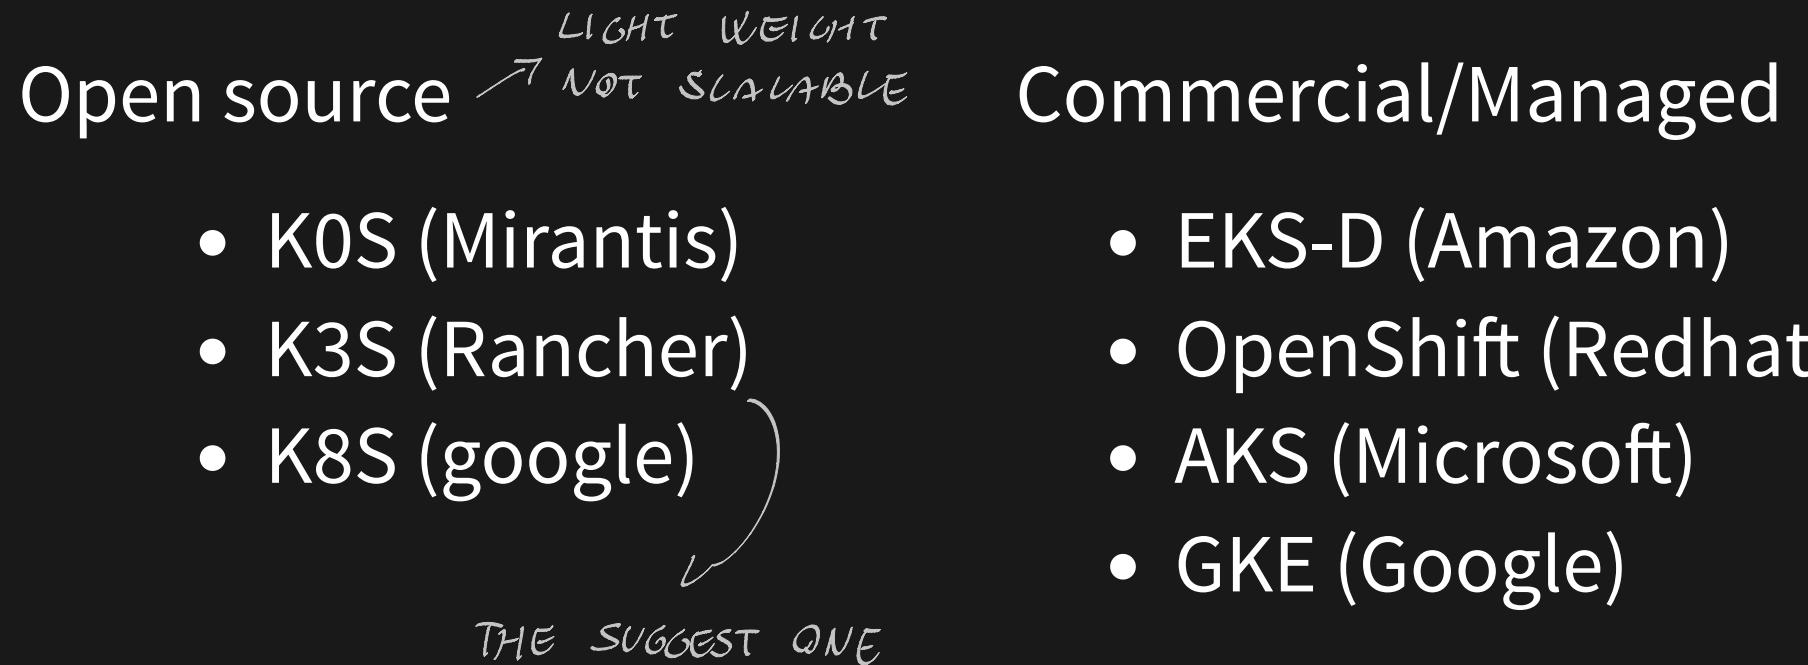
WHAT IS KUBERNETS?

- Project that was spun out of Google as an open source container orchestration platform.
- Built from the lessons learned in the experiences of developing and running Google's Borg.
- Designed from the ground-up as a loosely coupled collection of components centered around deploying, maintaining and scaling workloads.

IT WAS A DIFFERENT ENVIRONMENT BUT NOT OPEN SOURCE

CAN I JUST INSTALL KUBERNETES?

yes, and no, you need to choose a distribution before



WHO MANAGES KUBERNETES?

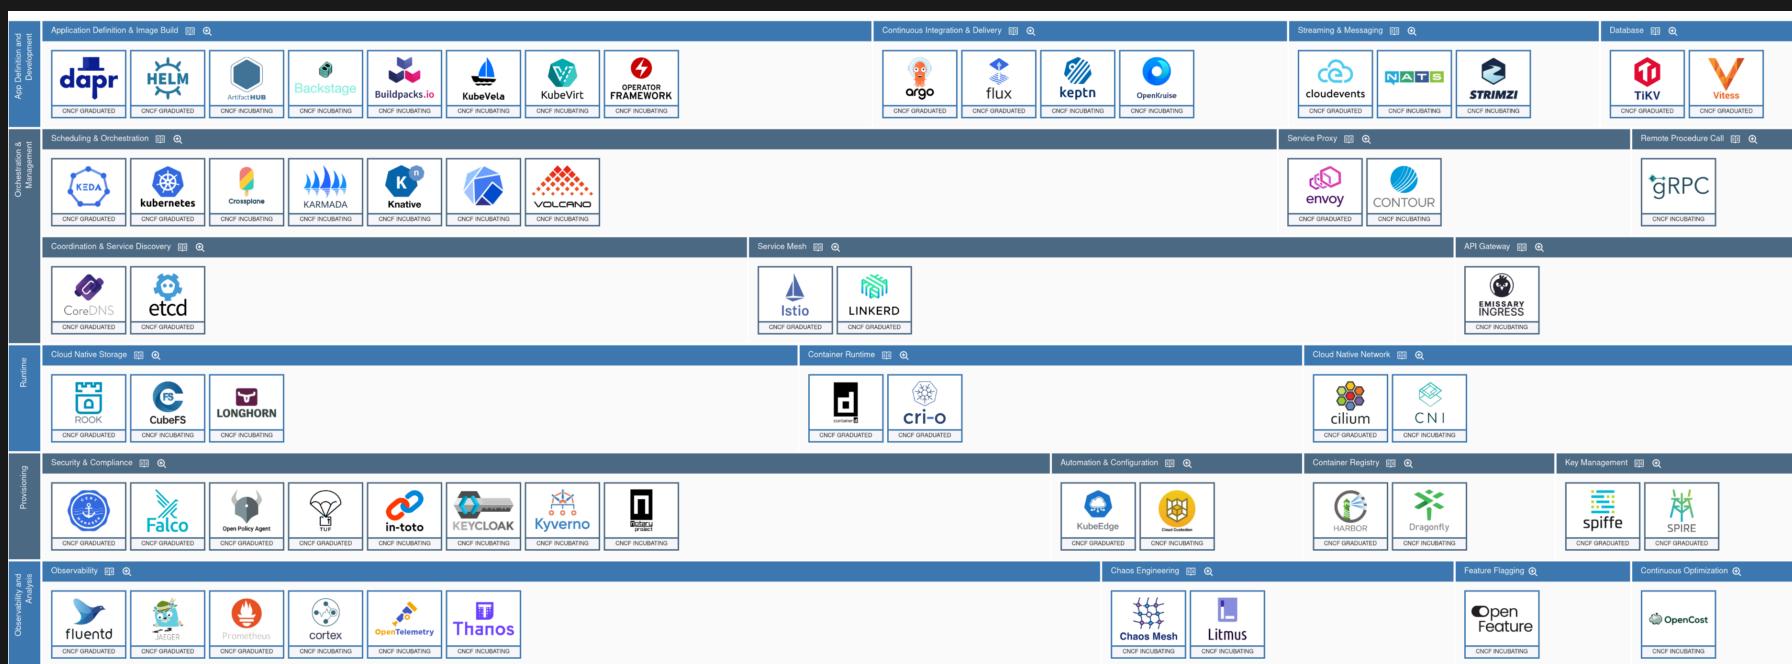


**CLOUD NATIVE
COMPUTING FOUNDATION**

The CNCF is a child entity of the Linux Foundation and operates as a vendor neutral governance group.

CNCF PROJECTS

(GRADUATED & INCUBATING)



WHAT DOES KUBERNETES PROVIDES YOU WITH?

- Service discovery and load balancing ex: DNS
- Storage orchestration
- Automated rollouts
- Automatic bin packing
- Self-healing
- Config management
- Batch execution
- Horizontal scaling
- IPv4/IPv6 dual-stack
- Extensibility

WHAT DOES KUBERNETES PROVIDES YOU WITH?

- Service discovery
- Storage orchestration
- Automated rollouts
- Automatic bin packing
- Self-healing
- Config management
- Batch execution
- Horizontal scaling
- IPv4/IPv6 dual-stack
- Extensibility

Ex. AUTOMATIC
MOUNTING FILE SYSTEM (Ex: NFS)
THAT CONTAINS ALL CONTAINERS ON
DIFFERENT MACHINE

WHAT DOES KUBERNETES PROVIDES YOU WITH?

- Service discovery
- Storage orchestration
- Automated rollouts and rollbacks
- Automatic bin packing
- Self-healing
- Config management
- Batch execution
- Horizontal scaling
- IPv4/IPv6 dual-stack
- Extensibility

MOVE FROM A VERSION
TO ANOTHER

WHAT DOES KUBERNETES PROVIDES YOU WITH?

- Service discovery
- Storage orchestration
- Automated rollouts
- Automatic process packing (process manager)
- Self-healing
- Config management
- Batch execution
- Horizontal scaling
- IPv4/IPv6 dual-stack
- Extensibility

WHAT DOES KUBERNETES PROVIDES YOU WITH?

- Service discovery
- Storage orchestration
- Automated rollouts
- Automatic bin packing
- Self-healing
- Config management
- Batch execution
- Horizontal scaling
- IPv4/IPv6 dual-stack
- Extensibility

↙
IF THE SYSTEM KEEP CRASHING
AND NEED TO UNDERSTAND WHAT
IS HAPPENING

WHAT DOES KUBERNETES PROVIDES YOU WITH?

- Service discovery
 - Storage orchestration
 - Automated rollouts
 - Automatic bin packing
 - Self-healing
 - Secret and configuration management
 - Batch execution
 - Horizontal scaling
 - IPv4/IPv6 dual-stack
 - Extensibility
- 
- KEEP
TRACK
OF CHANGE

WHAT DOES KUBERNETES PROVIDES YOU WITH?

- Service discovery
- Storage orchestration
- Automated rollouts
- Automatic bin packing
- Self-healing
- Config management
- Batch execution *→ MANAGE WORKLOAD*
- Horizontal scaling
- IPv4/IPv6 dual-stack
- Extensibility

WHAT DOES KUBERNETES PROVIDES YOU WITH?

- Service discovery
- Storage orchestration
- Automated rollouts
- Automatic bin packing
- Self-healing
- Config management
- Batch execution
- **Horizontal scaling**
- IPv4/IPv6 dual-stack
- Extensibility



CAN ADD MUCH CONTAINERS
AS YOU LIKE

WHAT DOES KUBERNETES PROVIDES YOU WITH?

- Service discovery
- Storage orchestration
- Automated rollouts
- Automatic bin packing
- Self-healing
- Config management
- Batch execution
- Horizontal scaling
- IPv4/IPv6 dual-stack
- Extensibility

WHAT DOES KUBERNETES PROVIDES YOU WITH?

- Service discovery
- Storage orchestration
- Automated rollouts
- Automatic bin packing
- Self-healing
- Config management
- Batch execution
- Horizontal scaling
- IPv4/IPv6 dual-stack
- **Designed for extensibility**

WHAT KUBERNETES DOES NOT

- limit the types of applications supported
- deploy source code/build your application
- provide application-level services
- dictate logging, monitoring, or alerting solutions.
- provide/mandate a configuration language
- comprehensive machine configuration,
maintenance, management, or self-healing systems.

YOU CAN HAVE 2 DIFFERENT KINDS OF APPLICATION

STATELESS VS STATEFUL

stateless: no past data nor state is stored or needs to be persistent when a new container is created (eg. CDN)

stateful: prior request history impacts the current state; the server must access and hold onto state information generated during the processing of the earlier request. (eg. DB)

WHAT KUBERNETES DOES NOT

- limit the types of applications supported
- **deploy source code/build your application** ↗^{> UNLESS YOU HAVE A PIPELINE}
- provide application-level services
- dictate logging, monitoring, or alerting solutions.
- provide/mandate a configuration language
- comprehensive machine configuration,
maintenance, management, or self-healing systems.

WHAT KUBERNETES DOES NOT

- limit the types of applications supported
- deploy source code/build your application
- provide application-level services *↗ IT DOES NOT GIVE MEMORY, BUT NEEDS A SOFTWARE*
- dictate logging, monitoring, or alerting solutions.
- provide/mandate a configuration language
- comprehensive machine configuration,
maintenance, management, or self-healing systems.

WHAT KUBERNETES DOES NOT

- limit the types of applications supported
- deploy source code/build your application
- provide application-level services *NEEDS A PLUG IN FOR LOG AND MONITORING*
- dictate logging, monitoring, or alerting solutions. *↑*
- provide/mandate a configuration language
- comprehensive machine configuration,
maintenance, management, or self-healing systems.

WHAT KUBERNETES DOES NOT

- limit the types of applications supported
- deploy source code/build your application
- provide application-level services
- dictate logging, monitoring, or alerting solutions.
- **provide/mandate a configuration language** *→ HIS OWN, BUT
HAVE MANY DEGREES OF
FREEDOM*
- comprehensive machine configuration,
maintenance, management, or self-healing systems.

WHAT KUBERNETES DOES NOT

- limit the types of applications supported
- deploy source code/build your application
- provide application-level services
- dictate logging, monitoring, or alerting solutions.
- provide/mandate a configuration language
- **comprehensive machine configuration,
maintenance, management, or self-healing systems.**

↳ IT DOESN'T CARE ABOUT THE MACHINE

WHAT DO YOU WANT WHEN YOU
EMBRACE KUBERNETES?

A DECLARATIVE INFRASTRUCTURE

WHAT IS A DECLARATIVE INFRASTRUCTURE?

Def: method of defining infrastructure in code by describing the desired state of resources.

Example Statement: "Deploy 3 instances of an application with load balancing enabled."

HOW IS THIS DIFFERENT FROM STANDARD MODE?

YOU JUST ASK AND DON'T NEED TO WRITE ALL NECESSARY CODE

Example:



- Me: “I want 3 healthy instances of redis to always be running.”
- Kubernetes: “Okay, I’ll ensure there are always 3 instances up and running.”

... something happens ...

- Kubernetes: “Oh look, one has died. I’m going to attempt to spin up a new one.”

Declarative

Specifies desired state

Eg.: YAML, JSON configurations

Idempotent

Less error-prone

Imperative

Specifies sequence of steps

Eg.: Shell scripts, Python

May cause issues on repeated runs

More complex troubleshooting

WHY IAC?

- Consistency: Enforces desired state consistently across environments.
- Scalability: Simplifies scaling
- Error Reduction: Eliminates complex scripting and minimizes human error.
- Auditability: Clear record of infrastructure states in version control.

A POSSIBLE FULL IAC STACK

- Terraform: Provisions the underlying infrastructure
VMs, networks, etc.
- Ansible: Configures the provisioned infrastructure
installing software, setting up configurations
- Kubernetes: Deploys/Manages containerized app.
on the provisioned & configured infrastructure.

OBVIOUSLY THE COMPLEXITY
RAISE A LOT

FROM DOCKER TO KUBERNETES

- Docker (single node)
- Docker compose (single node)
- Portainer, docker swarm (multi node)
Gray area, small medium environment
- Kubernetes (multi node orchestration)

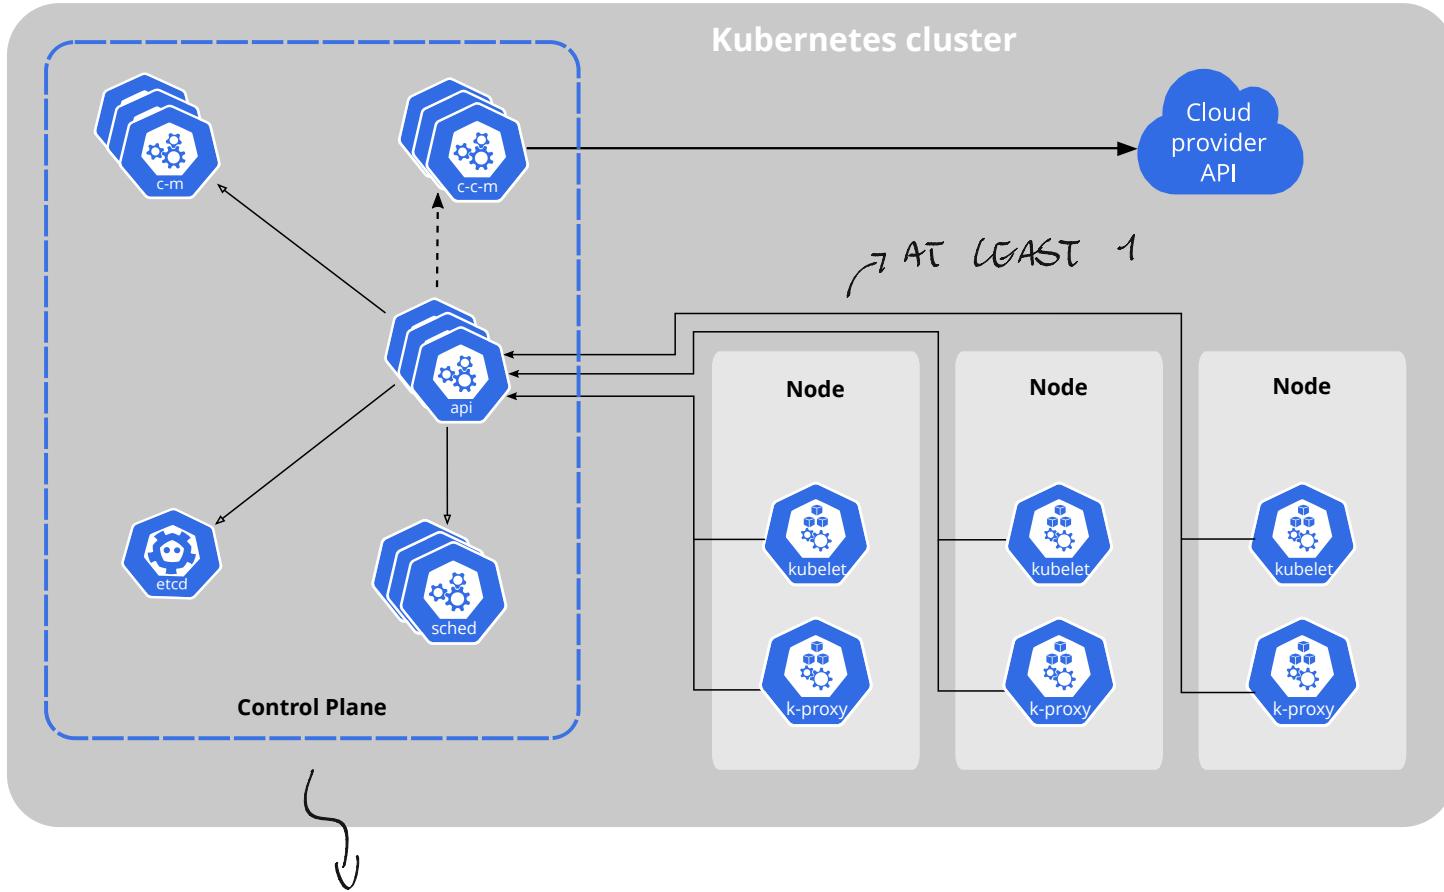
SHOULD I STILL USE DOCKER?

yes! Development and testing of your app can be done in docker.

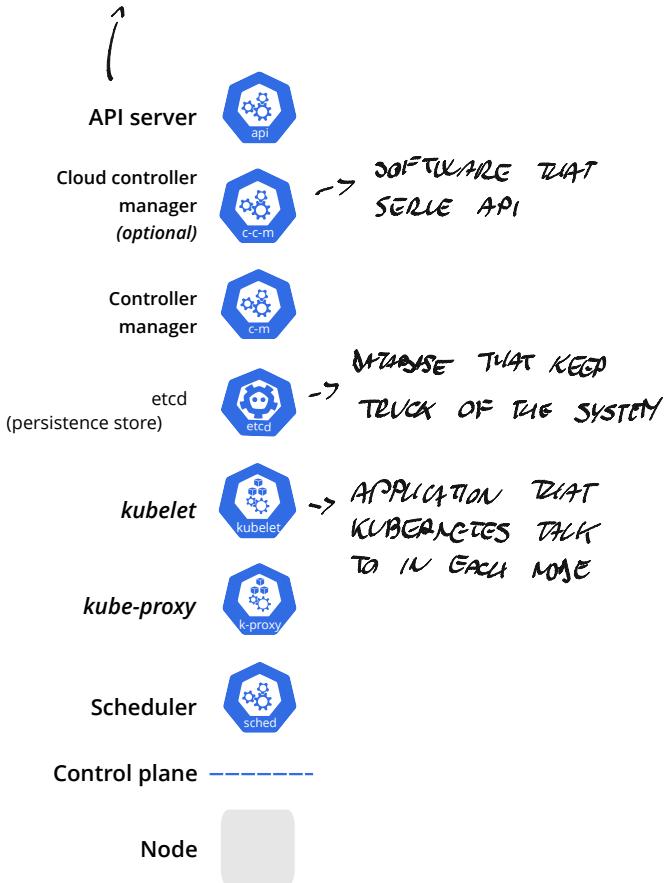
KUBERNETES IS NOT MADE FOR TESTING, DOCKER YES

COMPONENTS OVERVIEW

USUALLY YOU TALK WITH THE API
AND GET THE WORK DONE



SPECIAL NODE
WITH ALL INTERNAL
COMPONENTS



CONTROL PLANE COMPONENTS

- kube-apiserver
- etcd
- kube-scheduler
- kube-controller-manager
- cloud-controller-manager

exposes the Kubernetes API.

The API server is the front end for the Kubernetes control plane. The main implementation of a Kubernetes API server is kube-apiserver. kube-apiserver (main implementation) scale horizontally.

THE THING
YOU INTERACT
WITH

CONTROL PLANE COMPONENTS

- kube-apiserver
- etcd
- kube-scheduler
- kube-controller-manager
- cloud-controller-manager

key value store used as
backing store for all
cluster data.

CONTROL PLANE COMPONENTS

- kube-apiserver
- etcd
- **kube-scheduler**
- kube-controller-manager
- cloud-controller-manager

watches for newly created Pods with no assigned node, and selects a node for them to run on.

CONTROL PLANE COMPONENTS

- kube-apiserver
- etcd
- kube-scheduler
- **kube-controller-manager**
- cloud-controller-manager

scheduling decisions
based on:

- individual and collective resource
- hardware/software/policy constraints (ex: *DEPLOY A CONTAINER THAT NEED GPU ON A MACHINE THAT HAVE ONE*)
- affinity and anti-affinity
- data locality
- etc.

CONTROL PLANE COMPONENTS

- kube-apiserver
- etcd
- kube-scheduler
- **kube-controller-manager**
- cloud-controller-manager

runs controller processes:

- Node controller: Responsible for noticing and responding when nodes go down.
- Job controller: Watches for Job objects that represent one-off tasks, then creates Pods to run those tasks to completion.
- EndpointSlice controller: Populates EndpointSlice objects
- etc.

CONTROL PLANE COMPONENTS

- kube-apiserver
- etcd
- kube-scheduler
- kube-controller-manager
- **cloud-controller-manager**

cloud-specific control logic. only runs controllers that are specific to your cloud provider. If you are running Kubernetes on your own premises, there is no cloud controller manager.

NODE COMPONENTS

- kubelet
- kube-proxy (optional)
- Container runtime

An agent. The kubelet takes a set of PodSpecs and ensures that the containers described are running and healthy.

NODE COMPONENTS

- kubelet
- **kube-proxy (optional)**
- Container runtime

kube-proxy maintains network rules on nodes allowing network communication to your Pods from network sessions inside or outside of your cluster. CNI can have other proxy mechanisms.

NODE COMPONENTS

- kubelet
- kube-proxy (optional)
- Container runtime

It is responsible for managing the execution and lifecycle of containers within the Kubernetes environment. (container-d/CRIO)

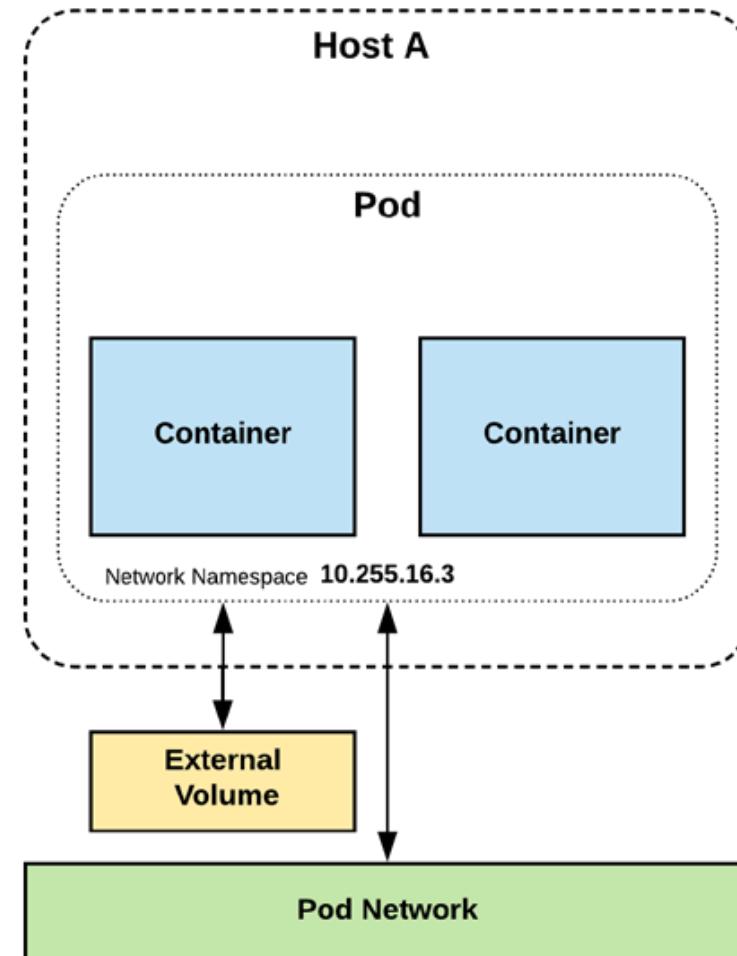
KUBE

SELECTED

OBJECTS

POD

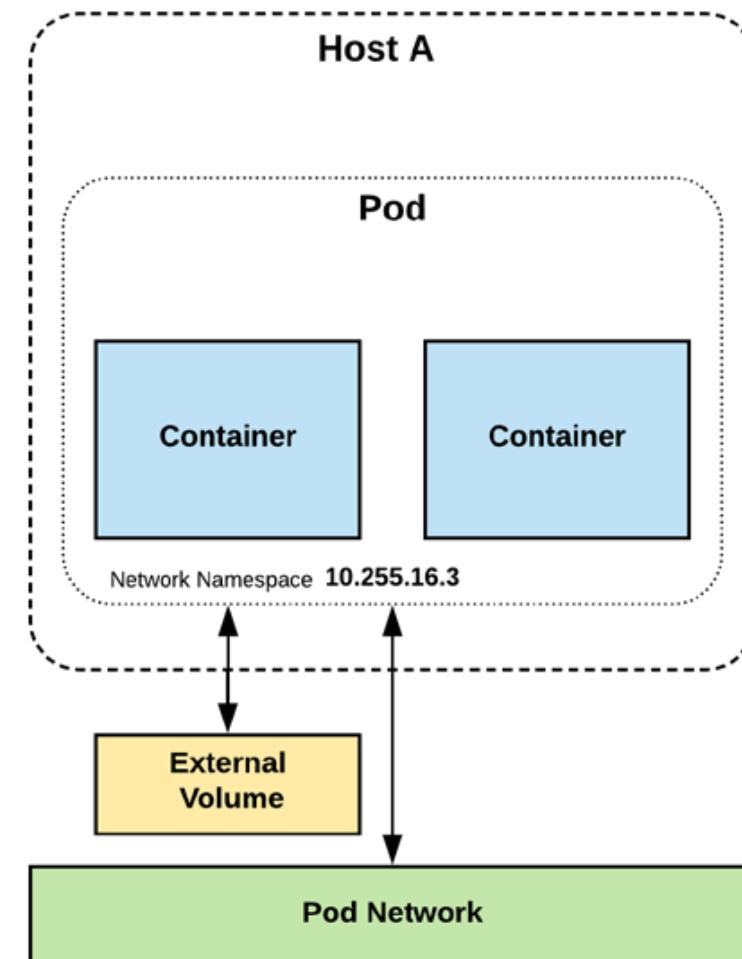
- smallest deployable units of computing
- Pods are one or MORE containers that share volumes, a network namespace, and are a part of a single context.div.
- pods are ephemeral



POD

You can use workload resources to create and manage multiple Pods:
for example:

- ReplicaSet -> Deployment
- StatefulSet
- DaemonSet



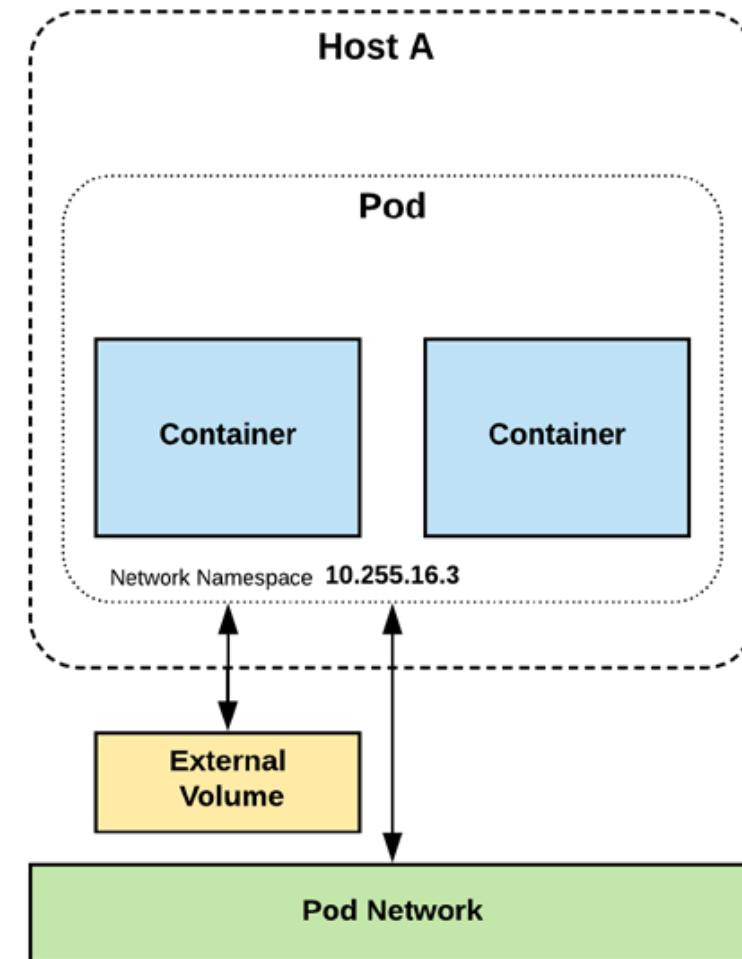
POD

- storage

A Pod can specify a set of shared storage volumes. All containers in the Pod can access the shared volumes

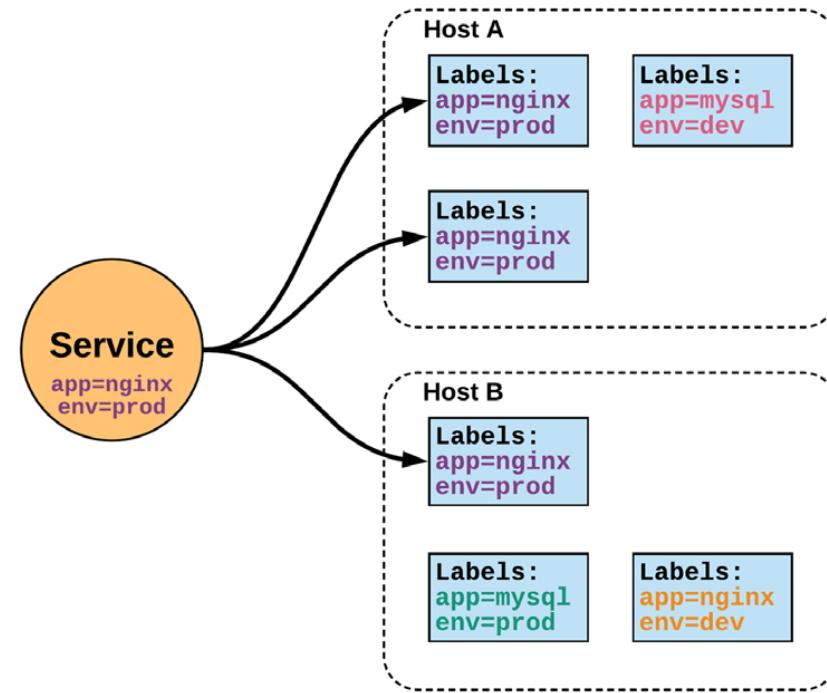
- network

Each Pod is assigned a unique IP address for each address family. Every container in a Pod shares the network namespace, including the IP address and network ports.



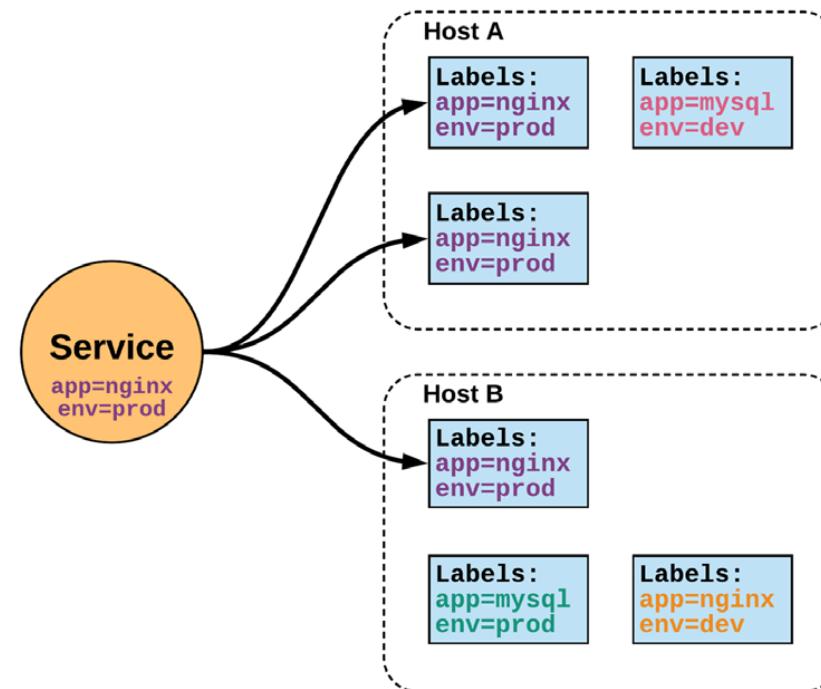
SERVICE

Expose an application running in your cluster behind a single outward-facing endpoint, even when the workload is split across multiple backends.



SERVICE

- Unified method of accessing the exposed workloads of Pods.
- Not ephemeral!



SERVICE

- Cluster IP
- Node port
- Load Balancer
- External Name

the Service on a cluster-internal IP.

SERVICE

- Cluster IP
- Node port
- Load Balancer
- External Name

the control plane allocates a port from a range. Each node proxies that port into your Service.

SERVICE

- Cluster IP
- Node port
- Load Balancer
- External Name

L2 loadbalancer, not available
without a plugin (like a metallb)

SERVICE

- Cluster IP
- Node port
- Load Balancer
- External Name

Maps the Service to the contents of the externalName field.

STORAGE

Data in the pod is not persistent (pods are ephemeral)

to ensure data persistency: persistent volume (PV)

PV? Represent physical storage resources, such as a block storage volume or a network share

STORAGE CLASSES:

- storage should be dynamically provisioned
- you need an external service for the storage Volume Provisioning

CONFIGMAP/SECRET

Are the correct way to store static config files/ secrets

INGRESS: THE GATEWAY TO SERVICES

L3 LOAD BALANCER

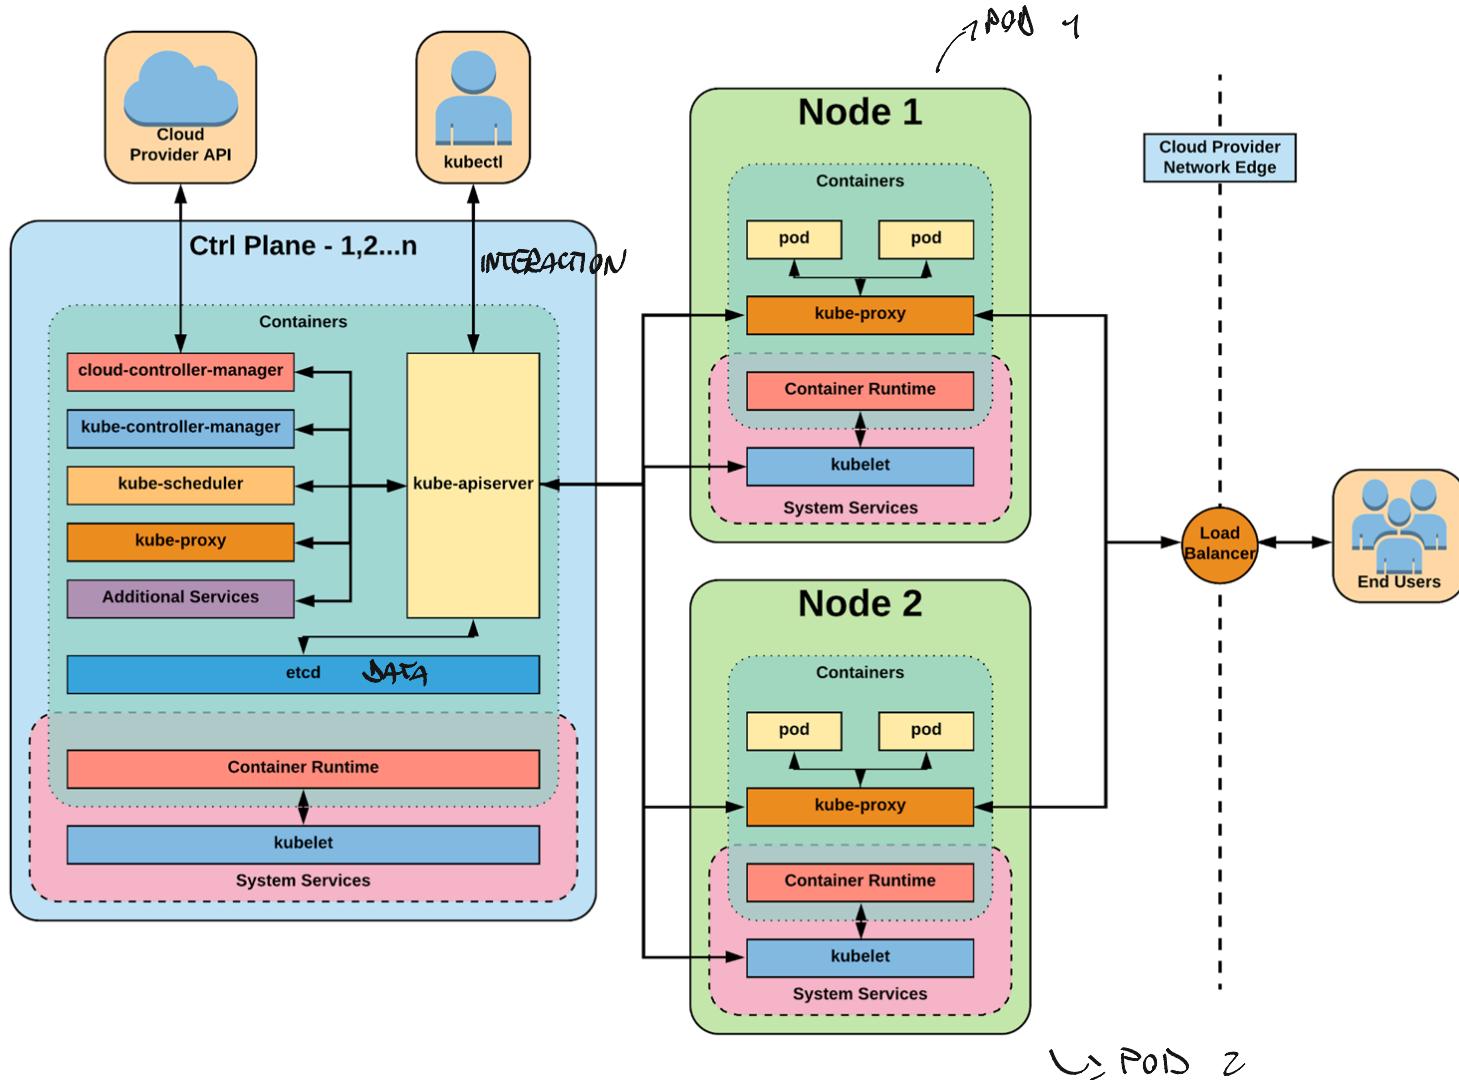
exposes HTTP/S (usually)services to the outside world

Acts as a reverse proxy, routing incoming traffic to different services based on URL paths, hostnames, or other criteria.

OBJECT EXAMPLE

```
1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   name: nginx-deployment
5 spec:
6   selector:
7     matchLabels:
8       app: nginx
9   replicas: 2
10  template:
11    metadata:
12      labels:
13        app: nginx
14    spec:
15      containers:
- NAME: NGINX
  IMAGE: NGINX:1.14.2
  PORTS:
- CONTAINERPORTS: 72
```

APPLICATION ARCHITECTURE



TAKEAWAY MESSAGES

WHAT IS KUBERNETES?

A powerful open-source platform for automating deployment, scaling, and management of containerized applications.

TAKEAWAY MESSAGES

KEY BENEFITS:

- Simplified Deployment
- Automated Scaling
- Self-Healing
- Load Balancing
- Container Orchestration

TAKEAWAY MESSAGES

CORE CONCEPTS (AT LEAST TWO):

- Pods
- Services